

Evaluating Program Semantics Reasoning with Type Inference in System F

Yifeng He, Luning Yang, Chris Gaw Gonzalo, Hao Chen



Introduction

Test-test compute (TTC) empowers coding LLMs the ability to reason on programs. However, are LLMs really “reasoning” on the semantics (logic) of the code?

- The evaluation gap for reasoning LLMs
 - Math: competition and answering problems
 - Code: Coding and patch
- How about reasoning about the logic behind math & code?
 - program-centric deductive system

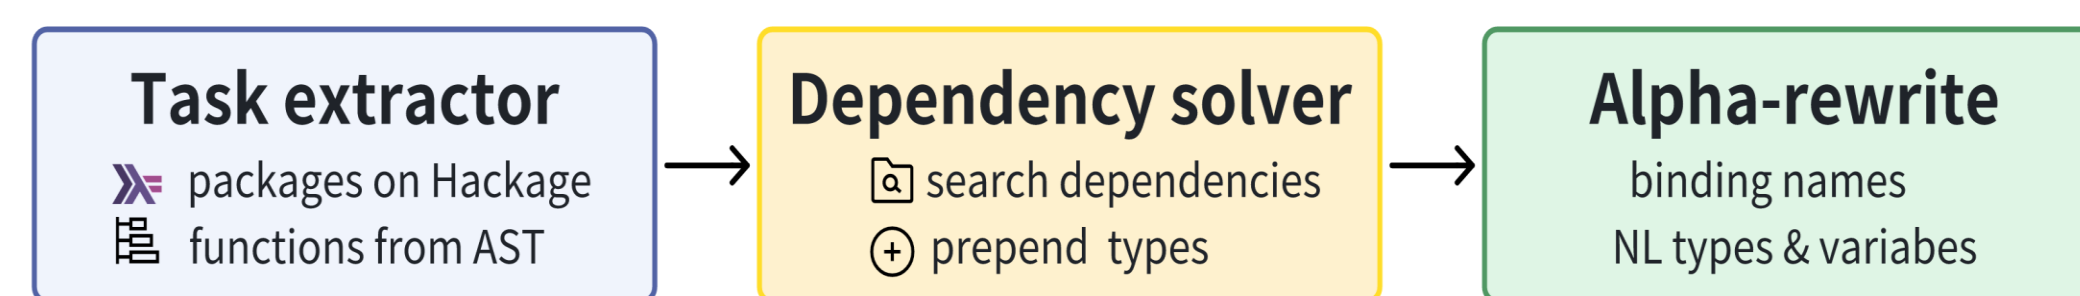
Background: propositions as types

- Types in programming \cong propositions in logic
 - Type inference is natural deduction
 - Type inference results are formally verifiable
 - Mutate the tasks, prove the results!
- \Rightarrow type inference is a task for *program semantics reasoning*

TF-Bench: benchmark construction

We use the Haskell Prelude to construct TF-Bench

- Formal deductive type system: System F (and System $F<$).
- The type signature is concise and is decoupled from the function body.
- Type equivalence is formally verifiable (w/ signature only).
- Haskell is the most popular language that meets these above conditions



Pipeline to construct TF-Bench.

Alpha-rewrite: removing natural language from tasks

We design three alpha-rewrite operators to transform the tasks:

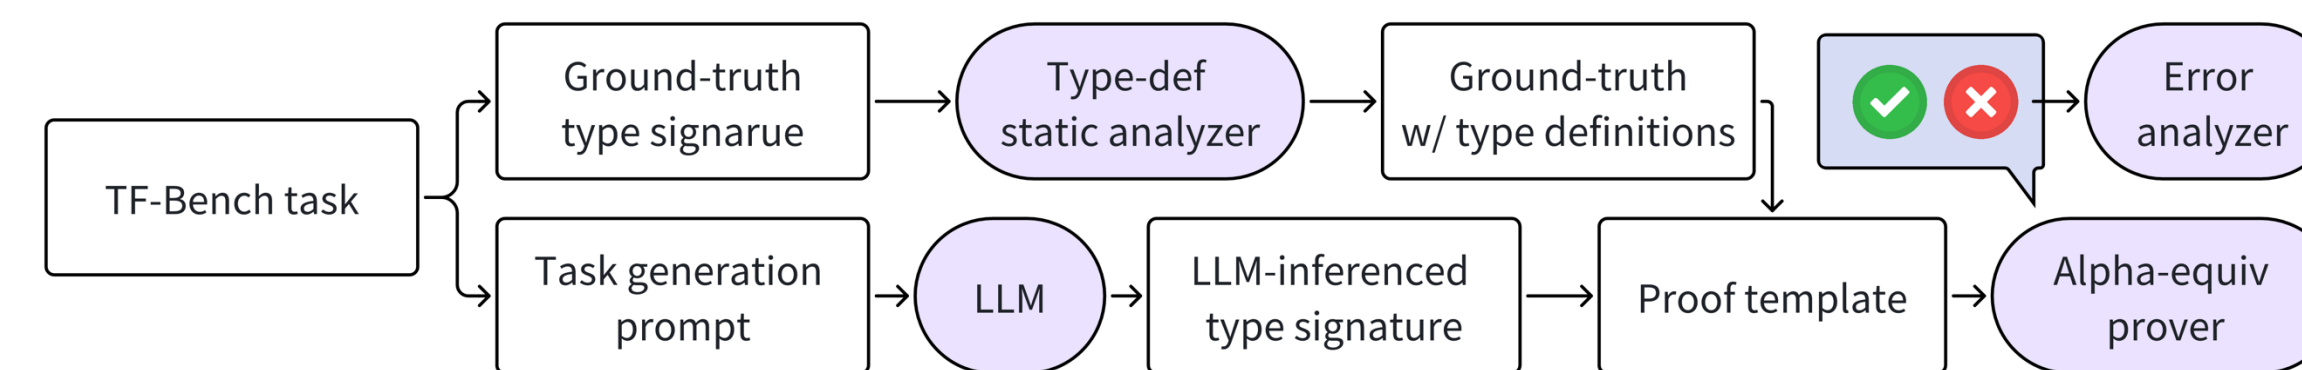
- NL-Type: `Int`, `Char`, `Bool`, `Eq`, `Ord` ... \rightarrow T1, T2, T3, T4, T5 ...
- Type-Var: `a`, `b`, `c`, `d`, `e` ... \rightarrow t1, t2, t3, t4, t5 ...
- Binding: `map`, `not`, `foldl`, `(+)` ... \rightarrow f1, f2, f3, f4 ...

Implementation details:

- Operators are *commutative* and *associative* under composition
- Transformed tasks *are alpha-equivalent* to the original tasks.

TF-Bench_{pure} = Alpha-rewrite(TF-Bench) reveals the true program semantics reasoning performance, w/o contamination & relying on NL-cues

Benchmark evaluation



Pipeline to evaluate LLMs on TF-Bench.

- Add definitions for new types after alpha-rewrite.
- Prompt the LLMs with tasks to generated a type signature.
- Construct a proof using ground-truth and generated type signatures.
- If the proof compiles, the two signatures are equivalent.
- Otherwise, we analyze what the error is.

Findings

- Reasoning robustness \Rightarrow LLMs (still!) can't reason on program semantics.
- Reasoning effectiveness \Rightarrow not all reinforcement learning are effective
- Fine-tuning on math data helps reasoning on code!



My homepage



Paper



Code

Results

RQ1. The performance gap

| Model | Version | TTC | Acc | Acc _{pure} | RS |
|-------------------|---------------|-----|-------|---------------------|-------|
| Claude-3.5-sonnet | 2024-06-20 | ✗ | 85.46 | 48.97 | 57.3 |
| Claude-3.7-sonnet | 2025-02-19 | ✓ | 90.42 | 55.85 | 61.77 |
| GPT-O3-mini | 2025-01-31 | ✓ | 90.43 | 48.40 | 53.52 |
| GPT-O3 | 2025-04-16 | ✓ | 81.91 | 52.66 | 64.29 |
| DeepSeek-V3 | 2025-03-25 | ✗ | 83.51 | 43.62 | 52.23 |
| DeepSeek-R1 | 2025-01-20 | ✓ | 86.70 | 44.15 | 50.92 |
| Qwen3 | 30B-A3B | ✓ | 81.38 | 40.43 | 49.68 |
| | 32B | ✓ | 87.94 | 43.09 | 49.00 |
| | 235B-A22B-FP8 | ✓ | 85.11 | 44.15 | 51.87 |

Table 1: Main evaluation results. $RS(m) = Acc_{pure}(m)/Acc(m)$.

RQ2. The effectiveness of TTC

| Model | TTC | Acc | Acc _{pure} | RE |
|-------------------|-----|-------|---------------------|------|
| Qwen3-235B-FP8 | ✗ | 80.49 | 35.64 | 1.37 |
| | ✓ | 86.70 | 44.15 | |
| Claude-3.7-sonnet | ✗ | 87.77 | 46.81 | 3.41 |
| | ✓ | 90.42 | 55.85 | |
| Gemini-2.5-flash | ✗ | 78.19 | 30.32 | 3.90 |
| | ✓ | 83.51 | 51.06 | |

Table 2: Reasoning effectiveness of top LLMs.

$$RE(m_{ttc}, m) = \frac{Acc_{pure}(m_{ttc}) - Acc_{pure}(m)}{Acc(m_{ttc}) - Acc(m)} = \frac{\Delta_{pure}}{\Delta}.$$

RQ3: Fine-tuning on math/code I

| FT Corpus | Base Model (FT Model) | Size | Acc | FT Acc | Δ | Acc _{pure} | FT Acc _{pure} | Δ_{pure} |
|-----------|-----------------------|------|-------|--------|----------|---------------------|------------------------|-----------------|
| Code | Gemma (CodeGemma) | 7B | 48.94 | 53.19 | + 4.25 | 7.45 | 12.23 | + 4.78 |
| | DeepSeek-V2 (-Coder) | 16B | 29.79 | 55.32 | + 25.53 | 7.98 | 15.96 | + 7.98 |
| | | 236B | 38.30 | 80.85 | + 42.55 | 11.17 | 36.70 | + 25.53 |
| | Mistral (Codestral) | 22B | 61.17 | 63.30 | + 2.13 | 19.68 | 11.17 | - 8.51 |
| | Qwen2.5 (-Coder) | 1.5B | 30.32 | 36.70 | + 6.38 | 6.91 | 9.04 | + 2.13 |
| | | 7B | 65.96 | 61.17 | - 4.79 | 21.28 | 21.28 | 0.00 |
| Math | Mistral (Mathstral) | 32B | 74.47 | 82.45 | + 7.98 | 36.17 | 31.91 | - 4.26 |
| | | 7B | 45.21 | 47.34 | + 2.13 | 7.99 | 15.43 | + 7.44 |
| | Qwen2 (-Math) | 7B | 40.43 | 43.09 | + 2.66 | 3.19 | 10.64 | + 7.45 |
| | | 72B | 63.83 | 71.28 | + 7.45 | 21.81 | 33.51 | + 11.7 |

Table 3: Result comparison of fine-tuning. FT Corpus: the corresponding fine-tuning corpus. Δ, Δ_{pure} : absolute increase in accuracy after fine-tuning.